

* syscall cont.

executing syscalls (b/c it's in a different addr space), is similar to remote procedure calls (RPCs).

```

////////////////////////////////////
/// verifying args

// example syscall: read(2)
int sys_read(int fd, __user void *buf, int size)
{
    // 1. verify args
    // fd:
    // - valid range? from 0 to some "MAX_FD"
    // - fd might not even be an opened file
    // - was file opened for reading?
    // - (does user process have permissions to read file? often checked at
    //   open time)

    // size: check not negative
    // - some OSs may have a max size of allowed read unit

    // buf (void* pointer -- belongs to user, virtual address!):
    // - not NULL
    // - is it allocated? -> is it in a valid addr space
    // kernel manages memory in 4KB (or PAGE_SIZE) units. Each page has
    // flags: Read, Write, eXecute. Protections are on the Virtual page addr.
    // Kernel needs to verify it can write to user addr.
    // Kernel translates the start addr of buf into the virt. page it belongs
    // to, then checks the virt-to-phys map for THIS process, to check if
    // there's even a mapping (else return error -EFAULT).
    // If there's a mapping, then verify access permissions to the page, else
    // get error (EFAULT, EPERM).
    // Kernel has to take buf, extend it to its enclosing 4K pages, then check
    // EACH of those pages for existence and permission.
    verify_area(buf, len, VERIFY_READ|VERIFY_WRITE);

    // 2. do actual work
    // - finds data from disk and file system, get I/O into memory
    // - kernel writes data from I/O buffer into user buffer,
    // - kernel cannot write directly to "__user void* buf"! Must write into
    // that buffer's physical location.

    // 3. cleanup
}

int sys_open(__user char *file, ...)
{
    char *kstr;
    // verify string: POSIX says PATHNAME_MAX is 4096
    // at most need to verify 2 pages.
    // then can translate virt to phys, read bytes to find null terminating
    // the string.
    // Best to copy string into kernel's own buffer. Like strdup(3).
    kstr = getname(__user char *str); // performs kmalloc, need to kfree
    // do putname when done with kstr.

    // can do it yourself using verify_area, then kmalloc, then copy from user
    // to kernel using:
    copy_from_user(void *kaddr, __user void *uaddr, int len);
    // or
    copy_to_user(__user void*u addr, void *kaddr, int len);
}

```

```

////////////////////////////////////
/// CODING STYLE

```

```

// ex. syscall needs to open 2 files to read, open 3rd file to write, and
// malloc a buffer.
// too much nesting, hard to code, fragile, bug prone, duplicates too much code
int sys_foo1(char *f1, char *f2, char *f3)

```

```

{
    void *kbuf;
    struct file *fp1, *fp2, *fp3;

    // 1. verify args

    // 2. prepare and initialize
    fp1 = filp_open(f1, O_READ, ...);
    if (failed to open fp1) {
        return -ERRNO; // return appropriate error
    }
    fp2 = filp_open(f2, O_READ, ...);
    if (failed to open fp2) {
        filp_close(fp2);
        return -ERRNO; // return appropriate error
    }
    fp3 = filp_open(f3, O_WRITE, ...);
    if (failed to open fp3) {
        filp_close(fp1);
        filp_close(fp2);
        return -ERRNO; // return appropriate error
    }
    kbuf = kmalloc(4096);
    if (kbuf == NULL) {
        filp_close(fp1);
        filp_close(fp2);
        filp_close(fp3);
        return -ENOMEM;
    }

    // main body of code

    // cleanup

}

```

```

// v2: w/o too much nesting
// benefits: single point of exit, no deep nesting, no duplicated code
// easier to code, less bug prone
int sys_foo2(char *f1, char *f2, char *f3)

```

```

{
    void *kbuf;
    struct file *fp1, *fp2, *fp3;
    int retval; // maybe initialize to something if needed

    // 1. verify args

    // 2. prepare and initialize
    fp1 = filp_open(f1, O_READ, ...);
    if (failed to open fp1) {
        retval = -EPERM; // or appropriate error
        goto out;
    }
    fp2 = filp_open(f2, O_READ, ...);
    if (failed to open fp2) {
        retval = -ERRNO; // return appropriate error
    }

```

```
    goto out_close1;
}
fp3 = filp_open(f3, O_WRITE, ...);
if (failed to open fp3) {
    retval = -ERRNO; // return appropriate error
    goto out_close2;
}
kbuf = kmalloc(4096);
if (kbuf == NULL) {
    retval = -ENOMEM;
    goto out_close3;
}

// main body of code
// do what needs to be done, then if succeeded, fall through to next label
// (out_kfree). Just remember to set retval = 0 (or whatever notes success)
// if there's a failure inside MIDDLE of main body of code: goto
// out_kfree.

// cleanup
out_kfree:
    kfree(kbuf);
out_close3:
    flip_close(fp3);
out_close2:
    flip_close(fp2);
out_close1:
    flip_close(fp1);
out:
    return retval;
}
```