

* async queuing system, cont.

How many consumers to run?

- depends on system conditions, and the nature of the work consumer does
- rule of thumb: have as many as you have free CPU cores
- if fewer consumers than CPU cores, cores idle, can get more throughput
- if more consumers than CPU cores, scheduling issues, not really better throughput

Prod/consume rates and queue's eventual states:

$R(p)$: rate that producers create new work

$R(c)$: rate that consumers process queued work

N : fixed size of queue

$R(p) \gg R(c)$: eventually queue is full, regardless of length, and producers have to be throttled.

$R(p) \ll R(c)$: eventually queue is empty, consumers are idle (waiting)

$R(p) \approx R(c)$: (producer rate is just SLIGHTLY faster than consumer, by some epsilon). Eventually, queue will also be full, but take longer to get to that steady state.

$R(p) < R(c)$: (producer rate is just SLIGHTLY lower than consumer, by some epsilon). Queue also empty eventually.

$R(p) = R(c)$: producer and consumer rates match perfectly. Queue is never empty or full. Very hard to get to perfect equilibrium. OSs and dist. systems. try very hard to adjust rates and #consumers and #producers to get close.

E.g., in linux, the `pdflush` with two thresholds:

- producers -- user processes touching file data
- consumers -- `pdflush` kthreads to write dirty data to f/s

problem: in systems w/ different I/O devices, SSDs are faster than HDDs. Can't assume that consumer rates are fixed.

New system: Backing Device Information (BDI)

- measure rates of consumers (I/O rates) to each media and even each f/s
- dynamic load balancing and proportional sharing
- e.g., HDD is 3 times slower than SSD, so can adjust consumer rates -- 3 consumers for SSD devices for each consumer to an HDD device.